

▼ 1 | BACKGROUND

1.1 | Context

The CalCOFI data set represents the longest (1949-present) and most complete (more than 50,000 sampling stations) time series of oceanographic and larval fish data in the world. It includes abundance data on the larvae of over 250 species of fish; larval length frequency data and egg abundance data on key commercial species; and oceanographic and plankton data. The physical, chemical, and biological data collected at regular time and space intervals quickly became valuable for documenting climatic cycles in the California Current and a range of biological responses to them. CalCOFI research drew world attention to the biological response to the dramatic Pacific-warming event in 1957-58 and introduced the term "El Niño" into the scientific literature.

CalCOFI conducts quarterly cruises off southern & central California, collecting a suite of hydrographic and biological data on station and underway. Data collected at depths down to 500 m include: temperature, salinity, oxygen, phosphate, silicate, nitrate and nitrite, chlorophyll, transmissometer, PAR, C14 primary productivity, phytoplankton biodiversity, zooplankton biomass, and zooplankton biodiversity.

1.2 | Scope

We would like to know if a connection between Salinity and Temperature exists

2 | REFERENCE LINKS

Main Website: <https://calcofi.com/>

Bottle-Dataset: <https://new.data.calcofi.com/index.php/database/calcofi-database/bottle-field-descriptions>

Cast-Dataset: <https://new.data.calcofi.com/index.php/database/calcofi-database/cast-table-column-descriptions>


▼ 3 | EDA (EXPLORATION DATA ANALYSIS)

▼ 3.1 | Import Packages and Datasets

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import os
import seaborn as sns
import plotly.express as px
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.tree import DecisionTreeRegressor
from sklearn.tree import plot_tree
from sklearn import tree
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor
```

```
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this ve
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
bot=pd.read_csv('/content/bottle.csv',low_memory=False)#.to_csv('bottlecsv.csv')
cast=pd.read_csv('/content/cast.csv',low_memory=False)#.to_csv('castcsv.csv')
```

 imagine.png

```
bot.info()
```

```
26 P04q          451786 non-null float64
27 SiO3uM        354091 non-null float64
28 SiO3qu        510866 non-null float64
29 NO2uM         337576 non-null float64
30 NO2q          529474 non-null float64
31 NO3uM         337403 non-null float64
32 NO3q          529933 non-null float64
33 NH3uM         64962 non-null float64
34 NH3q          808299 non-null float64
35 C14As1        14432 non-null float64
36 C14A1p        12760 non-null float64
37 C14A1q        848605 non-null float64
38 C14As2        14414 non-null float64
39 C14A2p        12742 non-null float64
40 C14A2q        848623 non-null float64
41 DarkAs        22649 non-null float64
42 DarkAp        20457 non-null float64
43 DarkAq        840440 non-null float64
44 MeanAs        22650 non-null float64
45 MeanAp        20457 non-null float64
46 MeanAq        840439 non-null float64
47 IncTim        14437 non-null object
48 LightP        18651 non-null float64
49 R_Depth        864863 non-null float64
50 R_TEMP        853900 non-null float64
51 R_POTEMP       818816 non-null float64
52 R_SALINITY     817509 non-null float64
53 R_SIGMA        812007 non-null float64
54 R_SVA          812092 non-null float64
55 R_DYNHT        818206 non-null float64
56 R_O2           696201 non-null float64
57 R_O2Sat       666448 non-null float64
58 R_SI03        354099 non-null float64
59 R_P04          413325 non-null float64
60 R_NO3          337411 non-null float64
61 R_NO2          337584 non-null float64
62 R_NH4          64982 non-null float64
63 R_CHLA        225276 non-null float64
64 R_PHAE0       225275 non-null float64
65 R_PRES        864863 non-null int64
66 R_SAMP        122006 non-null float64
67 DIC1           1999 non-null float64
68 DIC2           224 non-null float64
69 TA1            2084 non-null float64
70 TA2            234 non-null float64
71 pH2            10 non-null float64
72 pH1            84 non-null float64
73 DIC Quality Comment 55 non-null object
dtypes: float64(65), int64(5), object(4)
memory usage: 488.3+ MB
```

It has been decided to reduce the dataset to the following features:

```
bottle=bot.iloc[:,[0,4,5,6,7,8]]
bottle
```

	Cst_Cnt	Depthm	T_degC	Salnty	O2ml_L	STheta
0	1	0	10.500	33.4400	NaN	25.64900
1	1	8	10.460	33.4400	NaN	25.65600
2	1	10	10.460	33.4370	NaN	25.65400
3	1	19	10.450	33.4200	NaN	25.64300
4	1	20	10.450	33.4210	NaN	25.64300
...
864858	34404	0	18.744	33.4083	5.805	23.87055
864859	34404	2	18.744	33.4083	5.805	23.87072
864860	34404	5	18.692	33.4150	5.796	23.88911
864861	34404	10	18.161	33.4062	5.816	24.01426
864862	34404	15	17.533	33.3880	5.774	24.15297

864863 rows × 6 columns

```
bottle.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    Cst_Cnt    864863 non-null   int64
1    Depthm    864863 non-null   int64
2    T_degC    853900 non-null   float64
3    Salnty    817509 non-null   float64
4    O2ml_L    696201 non-null   float64
5    STheta    812174 non-null   float64
dtypes: float64(4), int64(2)
memory usage: 39.6 MB

```

▼ 3.2 | NAN VALUES ANALYSIS


Nan Values Visualization

```

plt.rcParams["figure.figsize"] = (8,5)
sns.heatmap(bottle.isna(),cbar=False,cmap='viridis')

```

<AxesSubplot: >



Data cleaning

```

cols=bottle.columns

for colonna in cols:
    bottle=bottle[~bottle[colonna].isnull()]

bottle=bottle.reset_index(drop=True)
plt.rcParams["figure.figsize"] = (8,5)
sns.heatmap(bottle.isna(),cbar=False,cmap='viridis')

```

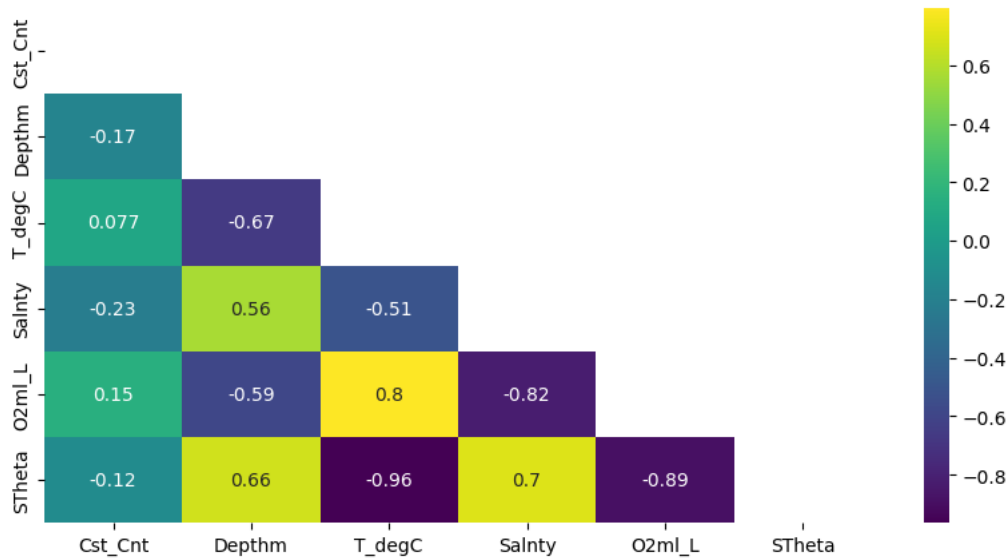
<AxesSubplot: >



3.3 | CORRELATION MATRIX



```
plt.figure(figsize=(10,5))
mask = np.triu(np.ones_like(bottle.corr(numeric_only = True), dtype=bool))
sns.heatmap(bottle.corr(numeric_only = True), cbar = True, annot = True, cmap="viridis", mask = mask);
```



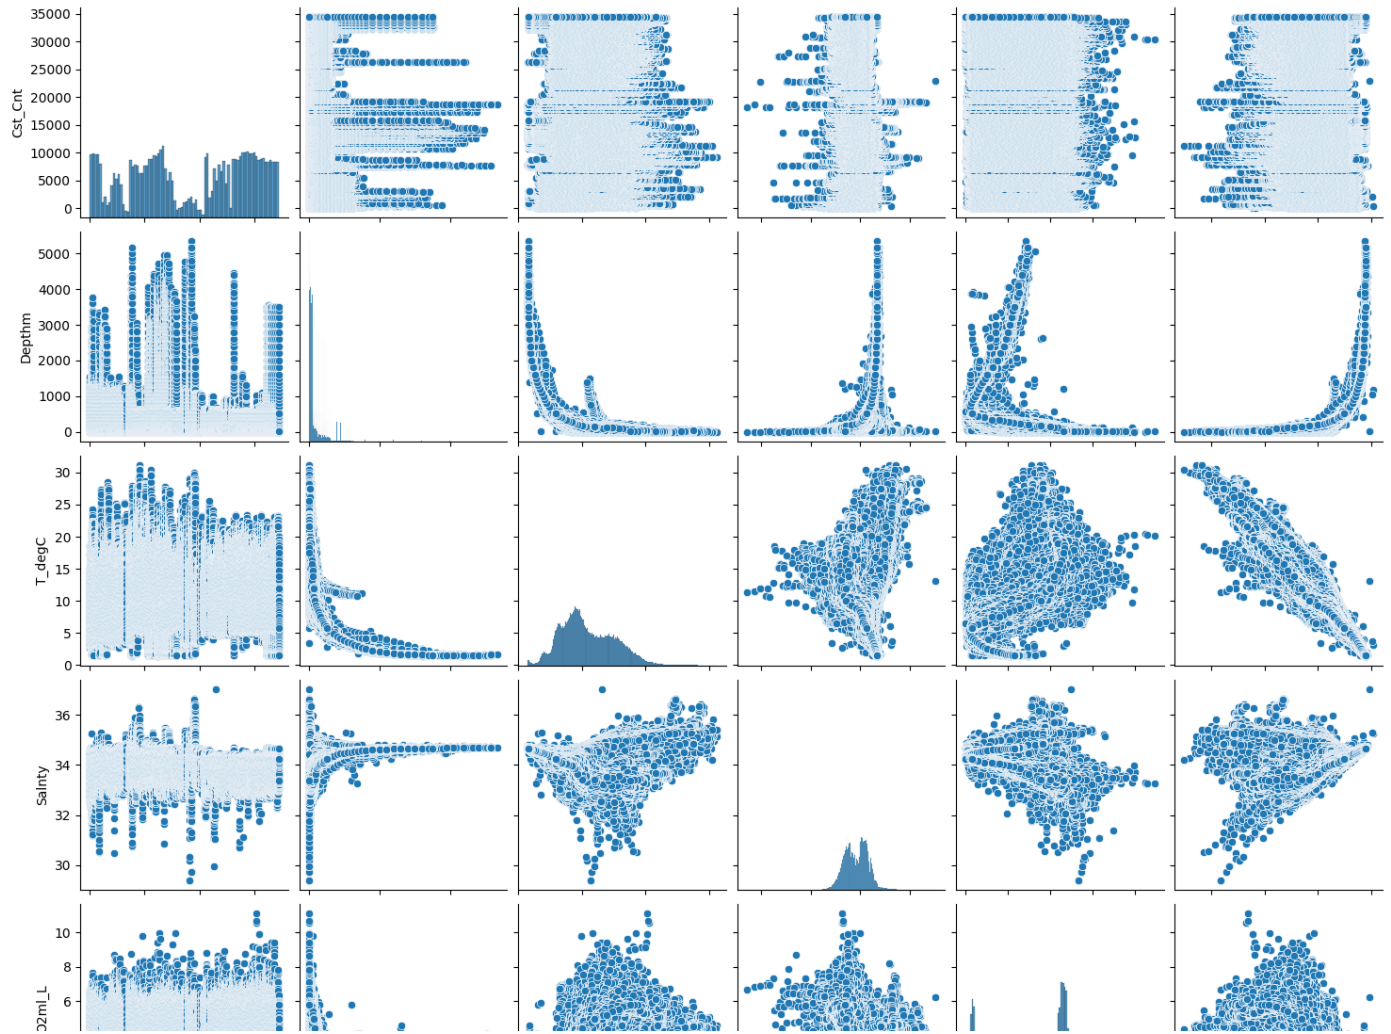
The Pearson's coefficient between Temperature and Salinity is not so high, -0.51

3.4 | FEATURES PLOTS


Plot pairwise relationships in the dataset

```
sns.pairplot(bottle)
```

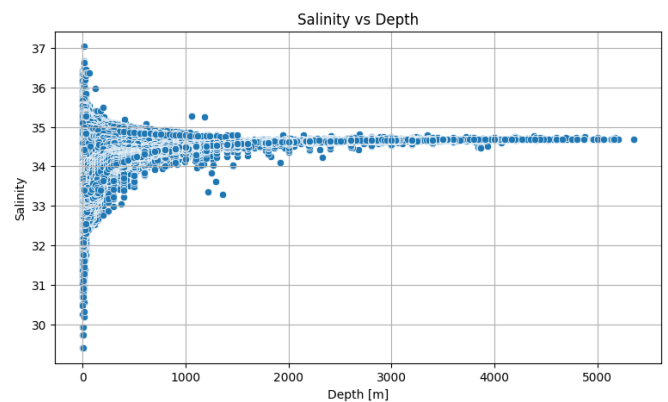

<seaborn.axisgrid.PairGrid at 0x7c9465acf7f0>



▼ 3.5 | DEPTH - TEMPERATURE - SALINITY



```
fig, ax = plt.subplots(1,2, figsize=(20,5))
g1=sns.scatterplot(data=bottle,x='Depthm',y='T_degC',ax=ax[0])
ax[0].set(xlabel='Depth [m]',ylabel='Temperature [°C]',title='Temperature vs Depth')
ax[0].grid()
g2=sns.scatterplot(data=bottle,x='Depthm',y='Salnty',ax=ax[1])
ax[1].set(xlabel='Depth [m]',ylabel='Salinity',title='Salinity vs Depth')
ax[1].grid()
```



As we can see from the image above temperature and salinity trace a hyperbolic trend as depth changes:

- In the first plot some points in the upper part of the image show a different trend from the rest of the graph. It's an unexpected behavior, because, for instance, at around 1000 m depth the expected temperature is around 5°C, but we also find temperatures between 10°C and 15°C (<https://www.windows2universe.org/earth/Water/temp.html>)
- In the second plot we can notice that the points converge asymptotically from both high and low salinity values. It means that in some areas of the ocean the salinity encreases with the depth, vice versa in other areas it decreases.

Distribution of Measuring Points on Geograph Datamap

Here we show the measuring points on the map using latitude and longitude informations.


```
df_merge=bottle.merge(cast, how='inner',on='Cst_Cnt')

fig = px.scatter_mapbox(df_merge,
                        lat="Lat_Dec",
                        lon="Lon_Dec",
                        #hover_name="Address",
                        hover_data=["Depthm"],
                        color="T_degC",
                        #color_continuous_scale=color_scale,
                        #size="Listed",
                        #animation_frame='Year',
                        zoom=3,
                        height=800,
                        width=800)

#fig.update_layout(mapbox_style="open-street-map")
fig.update_layout(
    mapbox_style="white-bg",
    mapbox_layers=[
        {
            "below": 'traces',
            "sourcetype": "raster",
            "sourceattribution": "United States Geological Survey",
            "source": [
                "https://basemap.nationalmap.gov/arcgis/rest/services/USGSImageryOnly/MapServer/tile/{z}/{y}/{x}"
            ]
        }
    ])
fig.show()
```

If we isolate points with salinity values in the range going from 34.85 to 38 we observe our alternative trend highlighted in the Temperature-Depth graph. It came out!

```
fig, ax = plt.subplots(1,2, figsize=(20,5))
bottle_filter=bottle[(bottle['Salnty']>34.85)&(bottle['Salnty']<38)]
g1=sns.scatterplot(data=bottle,x='Depthm',y='T_degC',ax=ax[0])
g1=sns.scatterplot(data=bottle_filter,x='Depthm',y='T_degC',ax=ax[0],legend=None,alpha=0.7)
ax[0].set(xlabel='Depth [m]',ylabel='Temperature [°C]',title='Temperature vs Depth')
ax[0].grid()
g2=sns.scatterplot(data=bottle,x='Depthm',y='Salnty',ax=ax[1])
g2=sns.scatterplot(data=bottle_filter,x='Depthm',y='Salnty',ax=ax[1],alpha=0.7)
ax[1].set(xlabel='Depth [m]',ylabel='Salinity',title='Salinity vs Depth')
ax[1].grid()
```



Distribution of Measuring Points for filtered dataset

Then we show the measuring points from the filtered dataset in the specific salinity range: 34.85 - 38.

```

df_merge_filter=bottle_filter.merge(cast, how='inner', on='Cst_Cnt')

fig = px.scatter_mapbox(df_merge_filter,
                        lat="Lat_Dec",
                        lon="Lon_Dec",
                        #hover_name="Address",
                        hover_data=["Depthm", "Salnty", "Cst_Cnt"],
                        color="T_degC",
                        #color_continuous_scale=color_scale,
                        #size="Listed",
                        #animation_frame='Year',
                        zoom=5,
                        height=800,
                        width=800)

#fig.update_layout(mapbox_style="open-street-map")
fig.update_layout(
    mapbox_style="white-bg",
    mapbox_layers=[
        {
            "below": 'traces',
            "sourcetype": "raster",
            "sourceattribution": "United States Geological Survey",
            "source": [
                "https://basemap.nationalmap.gov/arcgis/rest/services/USGSImageryOnly/MapServer/tile/{z}/{y}/{x}"
            ]
        }
    ])
fig.show()

```

We note that most points with high salinity values are located in the gulf.

▼ 3.6 | SALINITY - TEMPERATURE

```

plt.rcParams['figure.figsize']=10,6
ax=sns.scatterplot(data=bottle,x='Salnty',y='T_degC')
ax.set(xlabel='Salinity',ylabel='Temperature',title='Salinity vs Temperature')
plt.grid()

```

Salinity and Temperature don't seem to have a strong relationship.


We can try to use some additional information using the potential density (https://en.wikipedia.org/wiki/Potential_density). Remind that potential density is not a measured but a calculated quantity from Temperature, Pressure and Salinity.

Thanks to proportional link between chlorine and salinity, the last one (difficult to measure) is calculated through the measurement of chlorine, using an electrical conductivity meter.

The technology used to measure Conductivity, Temperature, and Pressure (related to Depth) at the same time in seawater is called CTD ([https://en.wikipedia.org/wiki/CTD_\(instrument\)](https://en.wikipedia.org/wiki/CTD_(instrument)))

Salinity-Temperature and potential Density plot

As predicted by the theoretical model, Temperature can be expressed as a function of Salinity and Density:


 imagine.png

We can try to highlight the density variable on the same graph:

```
theta_step=[20,23,24,25,26,27,34]
plt.rcParams['figure.figsize']=10,6

for theta1, theta2 in zip(theta_step, theta_step[1:]):
    bottle_stheta=bottle[(bottle['STheta']>=theta1) & (bottle['STheta']<theta2)]
    sns.scatterplot(data=bottle_stheta,y='T_degC',x='Salnty',label=str(theta1)+'-'+str(theta2))

plt.grid()
```



It works! But remind, it's not the real density at the specific depth, it's a potential density.

4 | PREDICTIVE MODELS

Split Train - Test

```
### Split
bottle_train, bottle_test = train_test_split(bottle, test_size=0.2, random_state=1234)
```



```

print("bottle_train: ", bottle_train.shape)
print("bottle_test: ", bottle_test.shape, "\n")

#features =['Salnty', 'STheta']
features =bottle.columns.drop('T_degC').tolist()
target = 'T_degC'
X_train = bottle_train[features].values
y_train = bottle_train[target].values
X_test = bottle_test[features].values
y_test = bottle_test[target].values

bottle_train: (529014, 6)
bottle_test: (132254, 6)

```

▼ 4.1 | LINEAR REGRESSION

Temperature from Salinity

```

features1=['Salnty'] #only Salinity feature
X_train1 = bottle_train[features1].values
y_train1 = bottle_train[target].values
X_test1 = bottle_test[features1].values
y_test1 = bottle_test[target].values

# creazione modello regressione lineare
model = LinearRegression()
model.fit(X_train1,y_train1)
y_predict1 = model.predict(X_test1)

```

Linear Regression Error

```

#calcolo errore modello di regressione
print(model)
mse_lin=mean_squared_error(y_test1, y_predict1)
RMSE_lin=mse_lin**(1/2.0)

print('RMSE_lin:', RMSE_lin)

LinearRegression()
RMSE_lin: 3.6326596110432914

```


Linear Regression Plot

```

g2=sns.regplot(x='Salnty',y="T_degC",data=bottle, color="red")
sns.scatterplot(data=bottle,x='Salnty',y='T_degC')
g2.set(xlabel='Salinity', ylabel='Temperature',title='Linear Regression')

```

```
[Text(0.5, 0, 'Salinity'),
Text(0, 0.5, 'Temperature'),
Text(0.5, 1.0, 'Linear Regression')]
```



It doesn't seem to exist a clear trend. It's consistently with the -0.5 Pearson's coefficient got from the confusion matrix. The linear regression is inversely proportional, but it doesn't seem a strong correlation. Of course we isolated the feature of Salinity from the others. We could try to use it with other physical quantities that we've maintained in this work

Multiple Linear Regression

```
# creazione modello regressione lineare
model = LinearRegression()
model.fit(X_train,y_train)
print(model)
y_predict2 = model.predict(X_test)

#calcolo errore modello di regressione
mse_multilin=mean_squared_error(y_test, y_predict2)
RMSE_multilin=mse_multilin**(1/2.0)

print('RMSE_multilin:',RMSE_multilin)
```


```
LinearRegression()
RMSE_multilin: 0.3323601966453793
```

Much better

4.2 | REGRESSION TREE

```
model_tree = DecisionTreeRegressor(random_state=44,max_depth=10)
model_tree.fit(X_train, y_train)
predictions = model_tree.predict(X_test)

plt.figure(figsize=(10,10))
#plot decision tree depth=4
tree.plot_tree(model_tree,feature_names=features,filled=True,max_depth=2);
```



Regression Tree Error

```
mse_tree = mean_squared_error(y_test, predictions)
RMSE_tree= mse_tree**(1/2)
#print("MSE: ", mse_tree)
print("RMSE_tree: ", RMSE_tree)
```

RMSE_tree: 0.18035450438720613

STheta <= 23.836

STheta <= 24.968

STheta <= 26.096

Depth <= 670.5

4.3 | KNN

```
k=100
model = KNeighborsRegressor(n_neighbors=k)
print(model)

model.fit(X_train,y_train)

pred_y = model.predict(X_test)

KNeighborsRegressor(n_neighbors=100)
```

KNN Error

```
print(model)
mse_knn = mean_squared_error(y_test, pred_y)
RMSE_knn = mse_knn**(1/2)
print("RMSE_KNN: ", RMSE_knn)
```

KNeighborsRegressor(n_neighbors=100)
RMSE_KNN: 1.3157370402918163

KNN algorithm is worse than multilinear regressor and decision tree regressor.

4.4 | FEATURE IMPORTANCE ANALYSIS

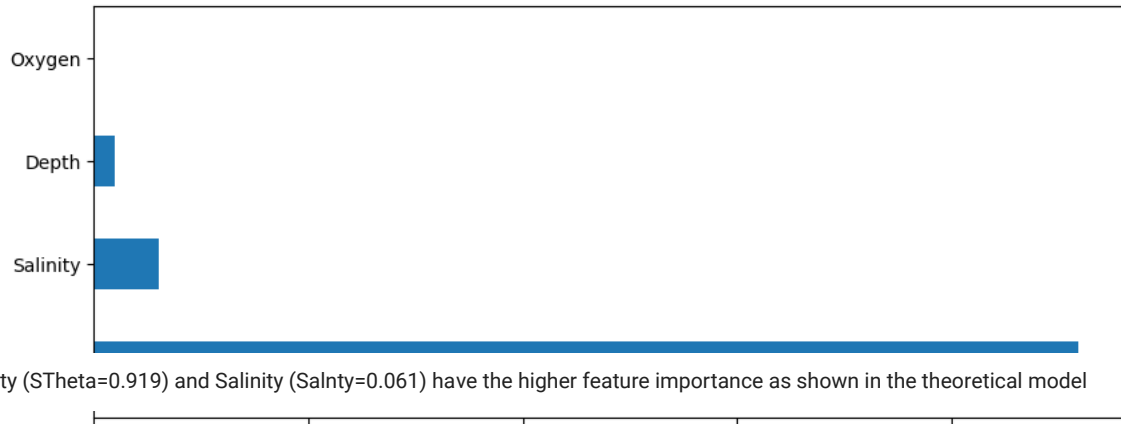
We can get a feature importance analysis through the decision tree model used above


```
for i,feature in enumerate(features):
    print(' ',feature,'=', model_tree.feature_importances_[i])
```

Cst_Cnt = 8.900949273136677e-06
Depth = 0.019813619791390027
Salnty = 0.06121591882241469
O2ml_L = 4.917424923116525e-05
STheta = 0.9189123861876909

```
pd.Series(model_tree.feature_importances_[1:5] , index=['Depth','Salinity','Oxygen','Density']).nlargest(4).plot(kind='barh',figsize=(10,4))
```

<AxesSubplot: >



immagine.png

So, we can try to repeat our predictions using only these forementioned two features

```
features2=['Salnty','STheta'] #only Salinity and Density features
X_train2 = bottle_train[features2].values
y_train2 = bottle_train[target].values
X_test2 = bottle_test[features2].values
y_test2 = bottle_test[target].values
```

```
# MULTIPLE LINEAR REGRESSION
model = LinearRegression()
model.fit(X_train2,y_train2)
y_predict2 = model.predict(X_test2)
```

```
#print(model)
#calcolo errore modello di regressione
mse_multilin=mean_squared_error(y_test2, y_predict2)
RMSE_multilin_2f=mse_multilin**(1/2.0)
```

```
print('RMSE_multilin_2f:',RMSE_multilin_2f)
```

```
#DECISION TREE
model = DecisionTreeRegressor(random_state=44,max_depth=10)
model.fit(X_train2, y_train2)
#print(model)
predictions = model.predict(X_test2)
mse_tree = mean_squared_error(y_test2, predictions)
RMSE_tree_2f= mse_tree**(1/2)
#print("MSE: ", mse_tree)
print("RMSE_tree_2f: ", RMSE_tree_2f)
```

```
#KNN
k=100
model = KNeighborsRegressor(n_neighbors=k)
model.fit(X_train2,y_train2)
pred_y = model.predict(X_test2)
#print(model)
mse_knn = mean_squared_error(y_test2, pred_y)
RMSE_knn_2f = mse_knn**(1/2)
print("RMSE_KNN_2f: ", RMSE_knn_2f)
```

```
#linear regression with density
features3=['STheta'] #only Salinity and Density features
X_train3 = bottle_train[features3].values
y_train3 = bottle_train[target].values
X_test3 = bottle_test[features3].values
y_test3 = bottle_test[target].values
```

```
# LINEAR REGRESSION with density
model = LinearRegression()
model.fit(X_train3,y_train3)
y_predict3 = model.predict(X_test3)
```

```
#calcolo errore modello di regressione per densità
mse_lin_dens=mean_squared_error(y_test3, y_predict3)
```

```

RMSE_lin_dens=mse_lin_dens**(1/2.0)
print("RMSE_dens: ", RMSE_lin_dens)

RMSE_multilin_2f: 0.4813688121553125
RMSE_tree_2f: 0.1778336794244693
RMSE_KNN_2f: 0.04434452522132766
RMSE_dens: 1.1215289238624324

```


RMSE Plot

```

Metodo_g1 = ['MultiLin','Tree','KNN']
Metodo_g2 = ['MultiLin', 'Tree','KNN']
Metodo_g3 = ['Lin']
Metodo_g4 = ['Lin']
RMSE_g1 = [RMSE_multilin_2f, RMSE_tree_2f, RMSE_knn_2f ]
RMSE_g2 = [RMSE_multilin, RMSE_tree, RMSE_knn]
RMSE_g3 = [RMSE_lin]
RMSE_g4 = [RMSE_lin_dens]

#setup figure
fig, ax = plt.subplots(1,4, figsize=(20,5))
g1=sns.barplot(x=Metodo_g1, y=RMSE_g1,ax=ax[3])
g2=sns.barplot(x=Metodo_g2, y=RMSE_g2,ax=ax[2])
g3=sns.barplot(x=Metodo_g3, y=RMSE_g3,ax=ax[0])
g4=sns.barplot(x=Metodo_g4, y=RMSE_g4,ax=ax[1])
g1.set_title('Prediction with Density and Salinity')
g2.set_title('Prediction with 5 features')
g3.set_title('Prediction with Salinity only')
g4.set_title('Prediction with Density only')
g1.set_ylim(0, 3.7)
g1.set_ylabel('RMSE')
g2.set_ylim(0, 3.7)
g2.set_ylabel('RMSE')
g3.set_ylim(0, 3.7)
g3.set_ylabel('RMSE')
g4.set_ylim(0, 3.7)
g4.set_ylabel('RMSE')
#g2.set_title('Boxplot della Densità')
plt.show()

```



5 | CONCLUSION

1. Apparently there isn't a strong relationship between Temperature and Salinity as shown in the picture on the right ('Prediction with Salinity only')
2. In order to enhance the temperature prediction with the linear regression model we need to use more features (the picture in the middle): Salinity, Density, Oxygen Saturation, Depth
3. KNN model shows better results when we use the only features with a higher importance: Density and Salinity. The same features introduced from the theoretical model on Temperature-Salinity-Density

Last but not least...

Below some plots to give a graphic explanation of why by adding the density feature we are able to improve the error.

One feature as Salinity is not enough to find a link with Temperature. A simple straight line is not able to describe a trend of a dependent variable (Temperature) function of two independent variables (Density and Salinity).

Instead adding one or more degrees of freedoms (i.e. Density) we give useful informations to our predictive model.


Temperature - Salinity- Density

```
theta_step=[20,23,24,25,26,27,34]
fig = plt.figure(figsize=[10,12])
ax = fig.add_subplot(projection='3d')

# set axis labels
ax.set_xlabel('\nTemperature')
ax.set_ylabel('\nSalinity')
ax.set_zlabel('\nDensity', linespacing=3.4)

for theta1, theta2 in zip(theta_step, theta_step[1:]):
    bottle_stheta=bottle[(bottle['STheta']>=theta1) & (bottle['STheta']<theta2)]
    ax.scatter(bottle_stheta['T_degC'], bottle_stheta['Salnty'], bottle_stheta['STheta'].values.reshape(-1,1))

plt.grid()
plt.gca().invert_zaxis()
plt.gca().view_init(-150, 30)
```



```

features3=['Salnty','STheta']
X_train3 = bottle_train[features3].values
y_train3 = bottle_train[target].values

model = LinearRegression()
model.fit(X_train3,y_train3)

#interpolazione dei punti del piano
N=10
x_mesh = np.linspace(30, 37, N) # intervalli sull'asse della salinità
y_mesh = np.linspace(20, 30,N) # intervalli sull'asse della densità
xx_pred, yy_pred = np.meshgrid(x_mesh, y_mesh)
model_viz = np.array([xx_pred.flatten(), yy_pred.flatten()]).T

predicted = model.predict(model_viz) # Valori di temperatura

```

The result of the linear regression is no longer a straight line, but a plane

```

fig = plt.figure(figsize=(10,12))
ax = fig.add_subplot(projection='3d') #crea immagine 3D

# set axis labels
ax.set_xlabel('\nTemperature')
ax.set_ylabel('\nSalinity')
ax.set_zlabel('\nDensity', linespacing=3.4)

ax.plot_trisurf(predicted, xx_pred.flatten(), yy_pred.flatten(), linewidth=0,alpha=0.6) # plot piano di regressione
ax.set_xlim(0,30)
ax.set_ylim(30,37)
ax.set_zlim(21,28)

#Visualizzazione con view: -170,40
plt.gca().view_init(-170, 40)
plt.gca().invert_zaxis()

# plot punti del dataframe Bottle
for theta1, theta2 in zip(theta_step, theta_step[1:]):
    bottle_stheta=bottle[(bottle['STheta']>=theta1) & (bottle['STheta']<theta2)]
    ax.scatter(bottle_stheta['T_degC'], bottle_stheta['Salnty'], bottle_stheta['STheta'].values.reshape(-1,1))

```

```

fig = plt.figure(figsize=(10,12))
ax = fig.add_subplot(projection='3d')

# set axis labels
ax.set_xlabel('\nTemperature')
ax.set_ylabel('\nSalinity')
ax.set_zlabel('\nDensity', linespacing=3.4)

ax.plot_trisurf(predicted, xx_pred.flatten(), yy_pred.flatten(), linewidth=0,alpha=0.3) # plot piano di regressione
ax.set_xlim(0,30)
ax.set_ylim(30,37)
ax.set_zlim(21,28)

#Visualizzazione con view: -170,30
plt.gca().view_init(-170, 30)
plt.gca().invert_zaxis()
# plot punti del dataframe Bottle
for theta1, theta2 in zip(theta_step, theta_step[1:]):
    bottle_stheta=bottle[(bottle['STheta']>=theta1) & (bottle['STheta']<theta2)]
    ax.scatter(bottle_stheta['T_degC'], bottle_stheta['Salnty'], bottle_stheta['STheta'].values.reshape(-1,1))

```

